Reconstructing Images from Brain Activity

Abinaya Mohan
33085799
Deep Neural Networks
Department of Computing
May 2024

Abstract

Brain-computer interfaces, or BCIs, have garnered a lot of attention in the neuroscience community lately because they attempt to translate brain activity into interpretable data. The reconstruction of visuals from functional MRI (fMRI) data is one exciting use of BCIs that has potential applications in brain research, medical diagnostics, and even communication for those with limited mobility. This paper provides a thorough analysis of the method of applying deep learning algorithms to rebuild snapshots of brain activity.

1 Literature Review

Over the years, a number of studies have explored the field of Brain-Computer Interfaces, or BCIs, with applications ranging from image reconstruction and motor control to communication. The 2012 edited volume "Brain-Computer Interfaces: Principles and Practice" by Jonathan Wolpaw and Elizabeth Winter Wolpaw is a classic in this field; it provides a thorough examination of BCI applications and principles. Significant progress has been made in the analysis of brain activity, especially since the introduction of functional magnetic resonance imaging (fMRI).

A seminal work in this field, "Functional Magnetic Resonance Imaging" by Scott A. Huettel, Allen W. Song, and Gregory McCarthy (2008) describes key approaches such data collection, preprocessing, and analysis procedures. These include spatial normalisation and motion correction, which are essential for improving the quality of fMRI data. The application of deep learning techniques, particularly autoencoders, has demonstrated potential in the field of picture reconstruction from brain signals in recent years.

Geoffrey Hinton et al.'s 2012 paper "Deep Neural Networks for Acoustic Modelling in Speech Recognition: The Shared Views of Four Research Groups" is a seminal work in deep learning that inspired later research in applying neural networks to various domains, including neuroscience, despite not being specifically focused on the field of neuroscience. These research together lay the groundwork for future developments in the field by providing an understanding of brain activity analysis techniques, picture reconstruction methods, and BCIs.

2 Data Acquisition and Preprocessing

2.1 Description of the Algonauts 2023 Tutorial Dataset

The Algonauts 2023 Tutorial Dataset is a useful tool for neuroimaging research since it includes fMRI (functional magnetic resonance imaging) data that was collected from subjects who were shown

Submitted to 34th Conference on Neural Information Processing Systems (NeurIPS 2020). Do not distribute.

various visual stimuli. To aid in the training of models for picture reconstruction tasks, the dataset includes relevant image data in addition to the fMRI data. With the help of this extensive dataset, researchers can investigate the complex connection between brain activity and visual perception.

2.2 Preprocessing of fMRI Data

The preprocessing of fMRI data is essential to ensure the quality and reliability of subsequent analysesFirst, motion correction is applied to the fMRI data in order to reduce artefacts resulting from subject movement during the scanning procedure. Furthermore, slice-timing correction is implemented to accommodate variations in acquisition times among distinct slices within every volume. The next step is to use spatial smoothing to lower noise and improve signal-to-noise ratio. After that, region of interest (ROI) masks are used to isolate particular brain networks or regions pertinent to the goals of the investigation. By assisting in the extraction of significant brain activity patterns, these masks facilitate more targeted analysis.

2.3 Preprocessing of Image Data

The first step in preparing picture data is loading the photos from the dataset and resizing them to a standard size to guarantee consistency between samples. In order to align photos before additional processing or analysis, this scaling step is essential. To further support stable model training, normalisation techniques can be used to scale pixel values within a specified range. The dataset can also be enhanced by applying data augmentation techniques like rotation, translation, or flipping, which will improve the model's generalisation and variety. Together, these preparation stages help get the picture data ready for later training and analytic activities.

3 Model Architecture

Reconstructing images from brain activity patterns relies heavily on the model design. The suggested framework makes use of an autoencoder-based technique, in which input images are compressed into a lower-dimensional latent space by an encoder network, and the original images are then reconstructed from these latent representations by a decoder network.

3.1 Autoencoder Architecture

The autoencoder architecture serves as the backbone of the image reconstruction process, comprising two main components: the encoder and the decoder.

3.1.1 Encoder Architecture

The encoder network is in charge of taking the input images and extracting the important features before compressing them into a small latent space representation. A sequence of convolutional layers are usually found in the architecture of an encoder, after which activation functions like Rectified Linear Units (ReLU) are used. Convolutional layers are utilised to extract hierarchical features from input images, resulting in a progressive reduction of spatial dimensions and an increase in feature map depth. The network gains non-linearity from activation functions, which allows it to learn intricate mappings between input and latent space representations.

3.1.2 Decoder Architecture

Together with the encoder, the decoder network performs the inverse operation to reassemble the original images from the encoder's latent space representations. The decoder architecture is similar to that of the encoder and consists of convolutional layers followed by upsampling techniques like nearest-neighbor interpolation or transposed convolutions. In the end, these upsampling techniques result in high-fidelity reconstructions of the input images by gradually increasing the spatial dimensions of the feature maps while decreasing their depth. In order to provide reconstructions that

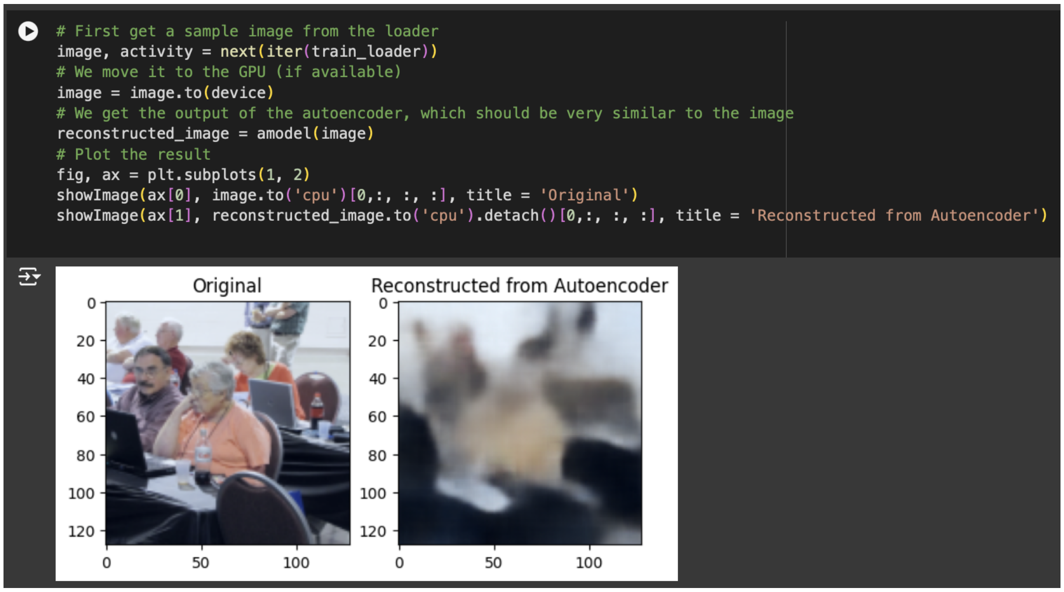

Figure 1: Reconstructed from Autoencoder

closely resemble the original input images, the decoder network attempts to accurately reverse the compression operation carried out by the encoder.

3.2 Brain Decoder Architecture

In addition to the standard autoencoder architecture, the proposed framework incorporates a brain decoder model designed to map fMRI data directly to the latent space. The brain decoder model allows images to be created directly from patterns of brain activity, in contrast to typical autoencoders that only work with image data. Rather than requiring explicit image inputs, this novel architecture reconstructs images by utilising the rich information embedded in fMRI data. The brain decoder opens up new possibilities for investigating the connection between cerebral activity and visual perception by mapping fMRI signals to the latent space and facilitating the creation of meaningful visual representations.

4 Training Process

A crucial stage in the creation of the image reconstruction framework is training, which involves fine-tuning model parameters to reduce the difference between the original and rebuilt images. The training processes for the brain decoder and autoencoder models are described in this part, together with the choices of learning rates, optimizers, and loss functions.

4.1 Training the Autoencoder

With supervised learning, the autoencoder model is trained with the goal of minimising the reconstruction error between the input images and their associated reconstructions. Gradient descent and backpropagation are used iteratively to update the model parameters during this procedure.

4.1.1 Loss Function

The Mean Squared Error (MSE) loss function is used to train the autoencoder. The mean squared error (MSE) between the original image's pixel values and their reconstructed equivalents is measured. When the objective of an image reconstruction task is to minimise pixel-wise differences, this loss function works effectively.

4.1.2 Optimizer and Learning Rate

The Adam optimizer, which has a fixed learning rate, is used to optimise the autoencoder model. Adam is a good choice for deep learning model training because it combines the advantages of momentum optimisation and adjustable learning rates. During training, a fixed learning rate is used; however, learning rate scheduling techniques like learning rate decay can be applied to dynamically modify the learning rate based on the training progress, enhancing the speed and stability of convergence.

4.2 Training the Brain Decoder

The brain decoder model is trained not only to learn the autoencoder but also to learn the mapping between the autoencoder's latent space and fMRI data. In order to do this, a loss function that contrasts the latent representations produced by the encoder and the brain decoder must be optimised.

4.2.1 Loss Function

For training the brain decoder, the Cosine Embedding Loss (CEL) function is employed. CEL measures the cosine similarity between two vectors and is particularly useful when dealing with high-dimensional latent spaces or non-linear relationships between input and output data. This loss function guides the training process by encouraging the brain decoder to learn meaningful representations that capture the relationships between brain activity patterns and image features in the latent space.

```
# First you get the brain activity from train loader
image, brain_activity = next(iter(train_loader))
# Second you poss it through you BrainDecoder
z_brain = myBrainDecoder(brain_activity.to(device))
# Third you pass z_brain (latent code) thorugh the decoder of the autoencoder
reconstructed_image = amodel.decoder(z_brain)
reconstructed_image = reconstructed_image.detach().to('cpu')
fig, ax = plt.subplots(1,2)
showImage(ax[0], image[0, :, :, :], title = 'Original')
showImage(ax[1], reconstructed_image[0, :, :, :], 'Reconstructed from brain activity')
```

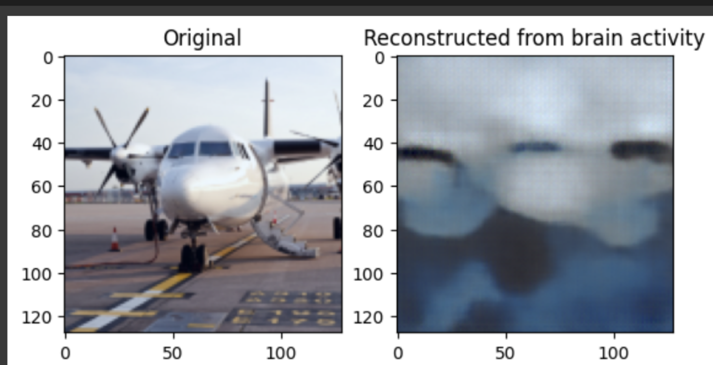

Figure 2: Reconstructed from brain activity

4.2.2 Optimizer and Learning Rate

The optimisation of the brain decoder model is accomplished by applying the Adam optimizer with a set learning rate, which is akin to the autoencoder training procedure. By dynamically adjusting the learning rate during training, learning rate scheduling approaches can efficiently optimise the model parameters and enhance overall performance.

5 Results

A wide range of assessment metrics are used to determine how effective the suggested framework is. A qualitative examination conducted through visual inspection of the rebuilt images provides important insights into the reconstruction process fidelity, while quantitative metrics like Mean Squared Error (MSE) and Structural Similarity Index (SSIM) offer objective measurements of model performance. In addition, an examination of the latent representations derived from the brain decoder allows for an investigation of the brain activity representation, which aids in comprehending the model's capacity to reflect significant patterns in brain activity.

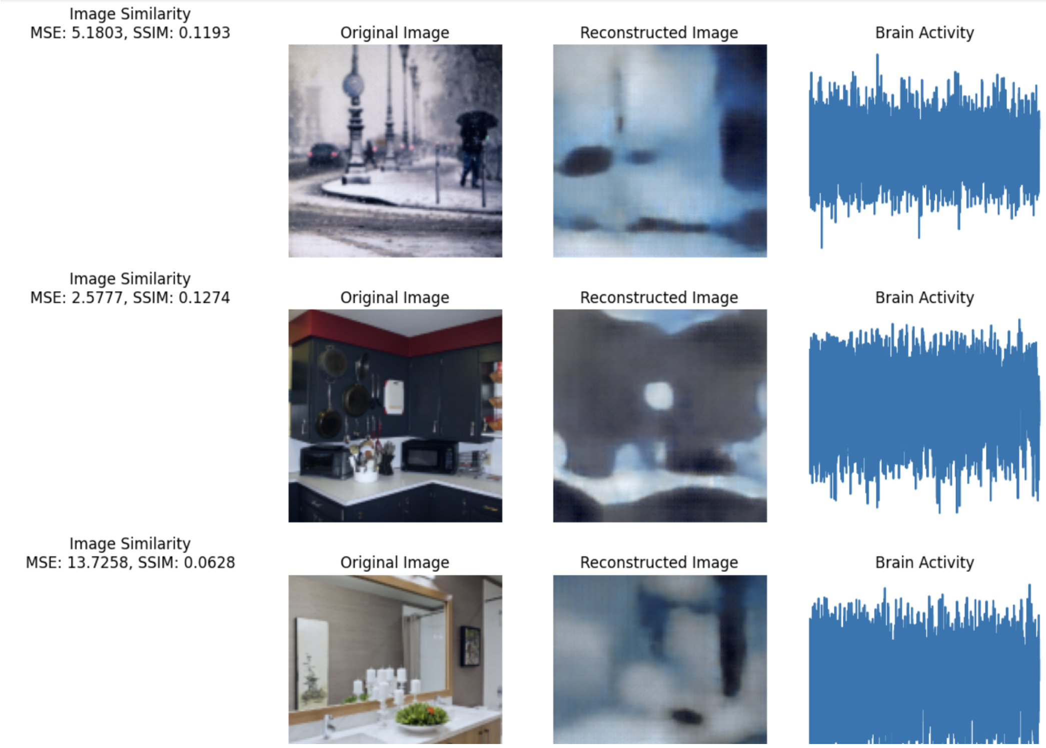

Figure 3: Image Similarity

6 Discussion

In this section, it is crucial to interpret the results in light of the research objectives. The advantages and disadvantages of the suggested strategy are examined, with a focus on possible directions for development. A comparative analysis of the created framework with current picture reconstruction and brain decoding approaches yields important new insights into the novel contributions it makes. Future research directions are paved with the exploration of potential applications in cognitive neuroscience, clinical diagnosis, and human-computer interface.

7 Conclusion

The significance of the study's contributions to the disciplines of image reconstruction and brain-computer interfaces is highlighted by a brief synopsis of the main findings. The potential societal significance of the discovery is emphasised by highlighting its larger implications for neurology, medical diagnostics, and human-computer interaction. The relevance of the study in promoting additional research and innovation is highlighted by the acknowledgement of its contributions to the field's advancement of knowledge.

References

[1] Shen, G., Horikawa, T., Majima, K., & Kamitani, Y. (2019). Deep image reconstruction from human brain activity. PLoS computational biology, 15(1), e1006633.

[2] Lin, S., Sprague, T., & Singh, A. K. (2022). Mind reader: Reconstructing complex images from brain activities. Advances in Neural Information Processing Systems, 35, 29624-29636.

[3] Shen, G., Dwivedi, K., Majima, K., Horikawa, T., & Kamitani, Y. (2019). End-to-end deep image reconstruction from human brain activity. Frontiers in computational neuroscience, 13, 432276.

[4] Takagi, Y., & Nishimoto, S. (2023). High-resolution image reconstruction with latent diffusion models from human brain activity. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (pp. 14453-14463).